



**Министерство науки и высшего образования Российской  
Федерации Федеральное государственное бюджетное  
образовательное учреждение высшего образования  
«Московский государственный технический университет  
имени Н.Э. Баумана (национальный исследовательский  
университет)» (МГТУ им. Н.Э. Баумана)**

**Работа №4**  
**по курсу «Технологии машинного обучения»**

**Выполнил**  
**студент группы ИУ5-64Б**  
**Шпак И.Д.**

**Москва, 2020**

# 1 Исходное задание

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую, с использованием `GridSearchCV` и/или `RandomizedSearchCV` и кросс-валидации, оцените качество оптимальной модели. Желательно использование нескольких стратегий кросс-валидации.
4. Обучите следующие модели:
  - одну из линейных моделей;
  - SVM;
  - дерево решений.
5. Оцените качество моделей с помощью двух подходящих для задачи метрик. Сравните качество полученных моделей.

# 2 Код программы

```
[105]: from IPython.display import Image
import numpy as np
import pandas as pd
from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.model_selection import KFold, RepeatedKFold, LeaveOneOut, LeavePOut,
↳ ShuffleSplit, StratifiedKFold
from sklearn.metrics import mean_absolute_error, mean_squared_error,
↳ mean_squared_log_error, median_absolute_error, r2_score,
↳ mean_absolute_percentage_error
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
```

```

from sklearn.linear_model import ElasticNet
from sklearn.svm import LinearSVR,SVR
from sklearn.preprocessing import OrdinalEncoder, StandardScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.datasets import load_wine
from io import StringIO
%matplotlib inline
pd.set_option("display.max_rows", None, "display.max_columns", None)
sns.set(style="ticks")

```

```

[106]: data = pd.read_csv("/home/igor/Downloads/CarPrice_Assignment.xls",sep=',')
data.shape

```

```

[106]: (205, 26)

```

```

[107]: cleanup_nums = {"doornumber":      {"four": 1, "two": 0},
                      "cylindernumber": {"four": 4, "six": 6, "five": 5, "eight": 8,
                                           "two": 2, "twelve": 12, "three": 3 },
                      "aspiration": {"std": 0, "turbo": 1},
                      "fueltype": {"gas": 0, "diesel": 1},
                      "enginelocation": {"front": 0, "rear": 1}}

data = data.replace(cleanup_nums)
data=pd.get_dummies(data, columns=["drivewheel"], prefix=["drive"])
data=pd.get_dummies(data, columns=["carbody"], prefix=["body"])
data["OHC_Code"] = np.where(data["enginetype"].str.contains("ohc"), 1, 0)
data.drop(data[(data['aspiration']=='turbo')].index,inplace=True)
data.drop(data[(data['fueltype']=='diesel')].index,inplace=True)
data.
↳ drop(["CarName","enginetype","fuelsystem","symboling","car_ID"],axis=1,inplace=True)

```

```

[108]: data_X = data.loc[:,data.columns]
clnm = StandardScaler()
data_X = clnm.fit_transform(data_X)
data_X = pd.DataFrame(data_X,columns=data.columns)
data_Y = data.loc[:, 'price']
data_X.drop(['price'],axis=1,inplace=True)
data_X_train, data_X_test, data_y_train, data_y_test = train_test_split(
    data_X, data_Y,test_size=0.2, random_state=360)
data_Y.head()

```

```
[108]: 0    13495.0
      1    16500.0
      2    16500.0
      3    13950.0
      4    17450.0
      Name: price, dtype: float64
```

```
[109]: data_X = data_X.to_numpy()
      data_Y = data_Y.to_numpy()
      data_X_train, data_X_test, data_Y_train, data_Y_test = train_test_split(
          data_X, data_Y, test_size=0.2, random_state=360)
```

```
[110]: data_X_train.shape
```

```
[110]: (164, 27)
```

```
[111]: reg = Ridge(alpha = 0.1).fit(data_X_train.reshape(-1, 27), data_Y_train)
      reg.coef_, reg.intercept_
      target1_0 = reg.predict(data_X_train)
      target1_1 = reg.predict(data_X_test)
      r2_score(data_Y_test, target1_1), mean_absolute_error(data_Y_test, target1_1)
```

```
[111]: (0.8866403748933834, 2181.032907732371)
```

```
[112]: scores = cross_val_score(Ridge(alpha = 1),
      data_X, data_Y,
      cv=4)
      print("%0.2f r^2 with a standard deviation of %0.2f" % (scores.mean(), scores.
      ↪std()))
```

```
0.34 r^2 with a standard deviation of 0.46
```

```
[113]: a = np.linspace(0.01, 1, 100)
      grid = GridSearchCV(estimator = Ridge(), param_grid={'alpha': a
      ↪a}, cv=RepeatedKfold(n_splits=3, n_repeats=3), scoring="r2")
      grid.fit(data_X, data_Y)
      grid.best_score_, grid.best_params_, grid.best_estimator_
```

```
[113]: (0.8481441870301406, {'alpha': 1.0}, Ridge())
```

```
[114]: grid.best_estimator_.fit(data_X_train, data_Y_train)
      target2_0 = grid.best_estimator_.predict(data_X_train)
      target2_1 = grid.best_estimator_.predict(data_X_test)
```

```
r2_score(data_Y_test, target2_1), mean_absolute_error(data_Y_test, target2_1)
```

```
[114]: (0.9102807327732504, 1937.558854553965)
```

```
[115]: scores = cross_val_score(grid.best_estimator_, data_X, data_Y,
    ↪cv=RepeatedKFold(n_splits=3, n_repeats=3))
print("%.2f r^2 with a standard deviation of %.2f" % (scores.mean(), scores.
    ↪std()))
```

0.82 r<sup>2</sup> with a standard deviation of 0.03

```
[116]: grid = GridSearchCV(estimator = Lasso(tol=1e-1) ,param_grid={'alpha': 
    ↪a},cv=RepeatedKFold(n_splits=3, n_repeats=3),scoring="r2")
grid.fit(data_X,data_Y)
grid.best_score_ , grid.best_params_,grid.best_estimator_
```

```
[116]: (0.8011735720066059, {'alpha': 1.0}, Lasso(tol=0.1))
```

```
[117]: grid.best_estimator_.fit(data_X_train, data_Y_train)
target3_0 = grid.best_estimator_.predict(data_X_train)
target3_1 = grid.best_estimator_.predict(data_X_test)
r2_score(data_Y_test, target3_1),1 mean_absolute_error(data_Y_test, target3_1)
```

```
[117]: (0.9231639328357497, 1792.2836196899398)
```

```
[118]: scores = cross_val_score(grid.best_estimator_, data_X, data_Y,
    ↪cv=RepeatedKFold(n_splits=3, n_repeats=3))
print("%.2f r^2 with a standard deviation of %.2f" % (scores.mean(), scores.
    ↪std()))
```

0.77 r<sup>2</sup> with a standard deviation of 0.12

```
[119]: b =np.linspace(0.1,1,10)
grid = GridSearchCV(estimator = ElasticNet(tol=1e-1) ,param_grid={'alpha': 
    ↪a,'l1_ratio' : b},cv=RepeatedKFold(n_splits=3, n_repeats=3),scoring="r2")
grid.fit(data_X,data_Y)
grid.best_score_ , grid.best_params_,grid.best_estimator_
```

```
[119]: (0.8199309484931719,
{'alpha': 0.8200000000000001, 'l1_ratio': 0.8},
ElasticNet(alpha=0.8200000000000001, l1_ratio=0.8, tol=0.1))
```

```
[120]: grid.best_estimator_.fit(data_X_train, data_Y_train)
target4_0 = grid.best_estimator_.predict(data_X_train)
target4_1 = grid.best_estimator_.predict(data_X_test)
r2_score(data_Y_test, target4_1), mean_absolute_error(data_Y_test, target4_1)
```

```
[120]: (0.9270503695691327, 1758.5738563522843)
```

```
[121]: scores = cross_val_score(grid.best_estimator_, data_X, data_Y,
    ↪cv=RepeatedKFold(n_splits=3, n_repeats=3))
print("%0.2f r^2 with a standard deviation of %0.2f" % (scores.mean(), scores.
    ↪std()))
```

```
0.81 r^2 with a standard deviation of 0.04
```

```
[122]: poly_model = Pipeline([('poly', PolynomialFeatures(degree=3)),
    ('linear', LinearRegression(fit_intercept=False))]
grid = GridSearchCV(estimator = poly_model ,param_grid={'poly__degree':
    ↪range(1,5,1)},cv=RepeatedKFold(n_splits=3, n_repeats=3),scoring="r2")
grid.fit(data_X,data_Y)
grid.best_score_ , grid.best_params_,grid.best_estimator_
```

```
[122]: (0.8239093877500444,
    {'poly__degree': 1},
    Pipeline(steps=[('poly', PolynomialFeatures(degree=1)),
    ('linear', LinearRegression(fit_intercept=False))]))
```

```
[123]: grid.best_estimator_.fit(data_X_train, data_Y_train)
target5_0 = grid.best_estimator_.predict(data_X_train)
target5_1 = grid.best_estimator_.predict(data_X_test)
r2_score(data_Y_test, target5_1), mean_absolute_error(data_y_test, target5_1)
```

```
[123]: (0.8817678394823523, 2225.7013621863985)
```

```
[124]: scores = cross_val_score(grid.best_estimator_, data_X, data_Y,
    ↪cv=RepeatedKFold(n_splits=3, n_repeats=3))
print("%0.2f r^2 with a standard deviation of %0.2f" % (scores.mean(), scores.
    ↪std()))
```

```
0.81 r^2 with a standard deviation of 0.03
```

```
[150]: reg1 = LinearSVR(C=1.0, loss='squared_epsilon_insensitive', max_iter=1000)
reg1.fit(data_X_train, data_Y_train)
target6_0=reg1.predict(data_X_train)
```

```
target6_1=reg1.predict(data_X_test)
r2_score(data_Y_test, target6_1), mean_absolute_error(data_y_test, target6_1)
```

[150]: (0.8994700594977598, 2049.1038225360794)

```
[126]: reg2 = DecisionTreeRegressor(random_state=360,max_depth=4)
reg2.fit(data_X_train, data_Y_train)
target7_0=reg2.predict(data_X_test)
sum(reg2.feature_importances_)
```

[126]: 1.0

```
[127]: r2_score(data_Y_test, target7_0),mean_absolute_error(data_Y_test, target7_0)
```

[127]: (0.9106421039901045, 1899.044949927358)

```
[128]: scores = cross_val_score(reg2, data_X, data_Y, cv=RepeatedKFold(n_splits=3,
↳n_repeats=3))
print("%0.2f r^2 with a standard deviation of %0.2f" % (scores.mean(), scores.
↳std()))
```

0.85 r<sup>2</sup> with a standard deviation of 0.05

```
[129]: #
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor,
↳export_graphviz
import graphviz
import pydotplus
data_X = data.loc[:,data.columns]
data_X.drop(['price'],axis=1,inplace=True)
def get_png_tree(tree_model_param, feature_names_param):
    dot_data = StringIO()
    export_graphviz(tree_model_param, out_file=dot_data,
↳feature_names=feature_names_param,
                    filled=True, rounded=True, special_characters=True)
    graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
    return graph.create_png()
Image(get_png_tree(reg2, data_X.columns), height='70%')
```

[129]:

